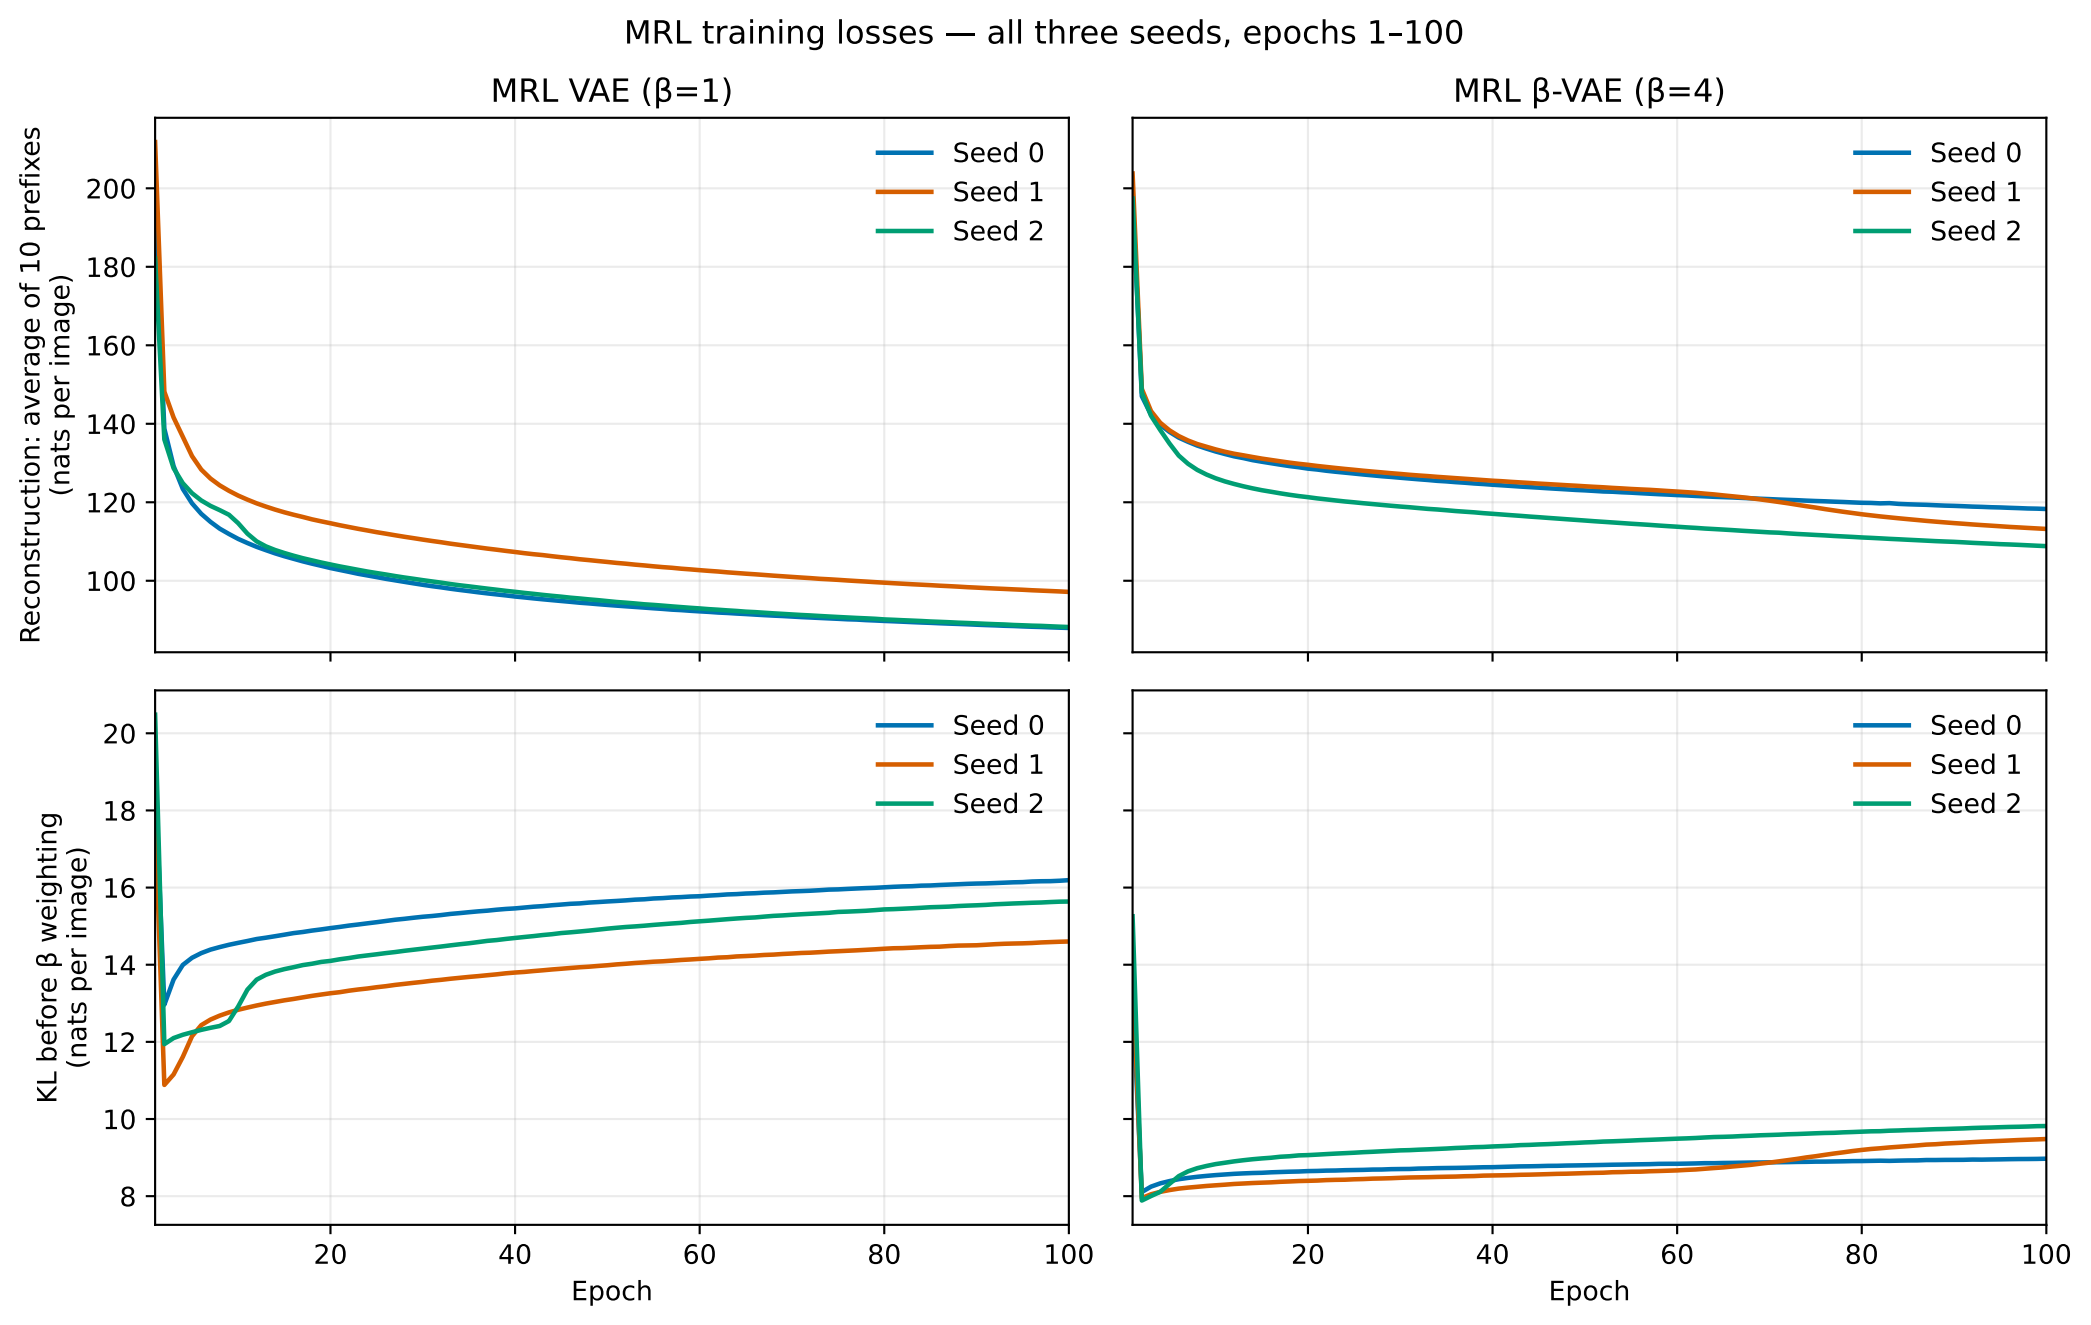
MRL training losses — all three seeds, epochs 1–100
MRL VAE (β=1)
MRL β-VAE (β=4)
Reconstruction: average of 10 prefixes (nats per image)
KL before β weighting (nats per image)
Seed 0
Seed 1
Seed 2
Epoch